

# Caching

*DS 5110/CS 5501: Big Data Systems*

*Spring 2024*

Lecture 2d

Yue Cheng



Some material taken/derived from:

- Wisconsin CS 544 by Tyler Caraza-Harter.

@ 2024 released for use under a [CC BY-SA](https://creativecommons.org/licenses/by-sa/4.0/) license.


# Learning objectives

- Describe the cache hierarchy
- Understand spatial locality and temporal locality
- Trace through access patterns with FIFO and LRU caching policies
  - Calculate cache performance metrics


# Outline

- Challenge: latency
- Cache hierarchy
  - CPU, RAM, SSD, Disk, Network
  - Tradeoffs
- Data access patterns, data locality, data access granularity
  - Spatial locality
  - Temporal locality
  - Cache lines and locality optimization
- What data should be cached?
  - Eviction policies: FIFO, LRU


# Interaction between CPU and RAM



# Interaction between CPU and RAM




# Load and store




**Challenge:** If we want to add some numbers stored in RAM, we need to **load** before adding and **store** after

# Latency to load from RAM



Very slow, but not long enough to switch to a different thread...

# Latency

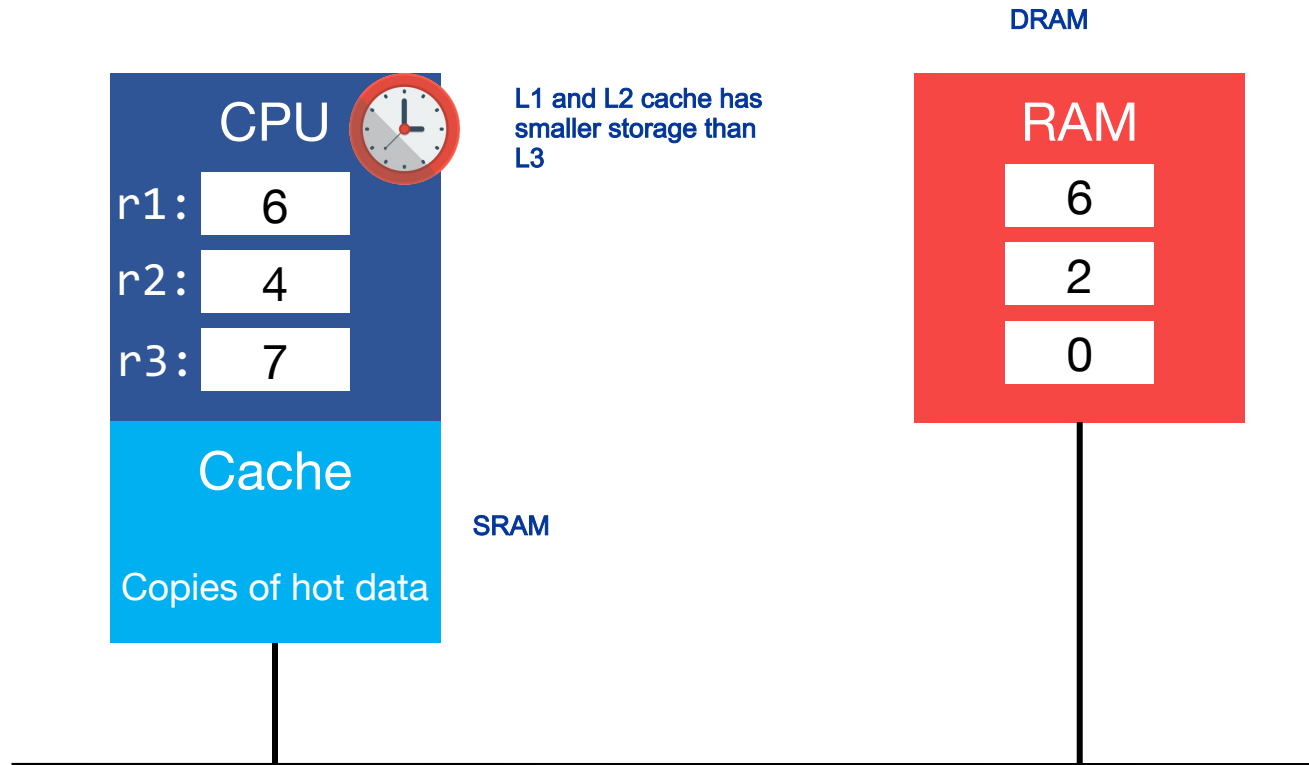


“How much time” is a **latency** measure.

**Throughput** (bytes/second) depends on how many loads we can do simultaneously.



# CPU Cache



**Idea:** CPUs can have a small but very fast memory built in for data that is frequently accessed

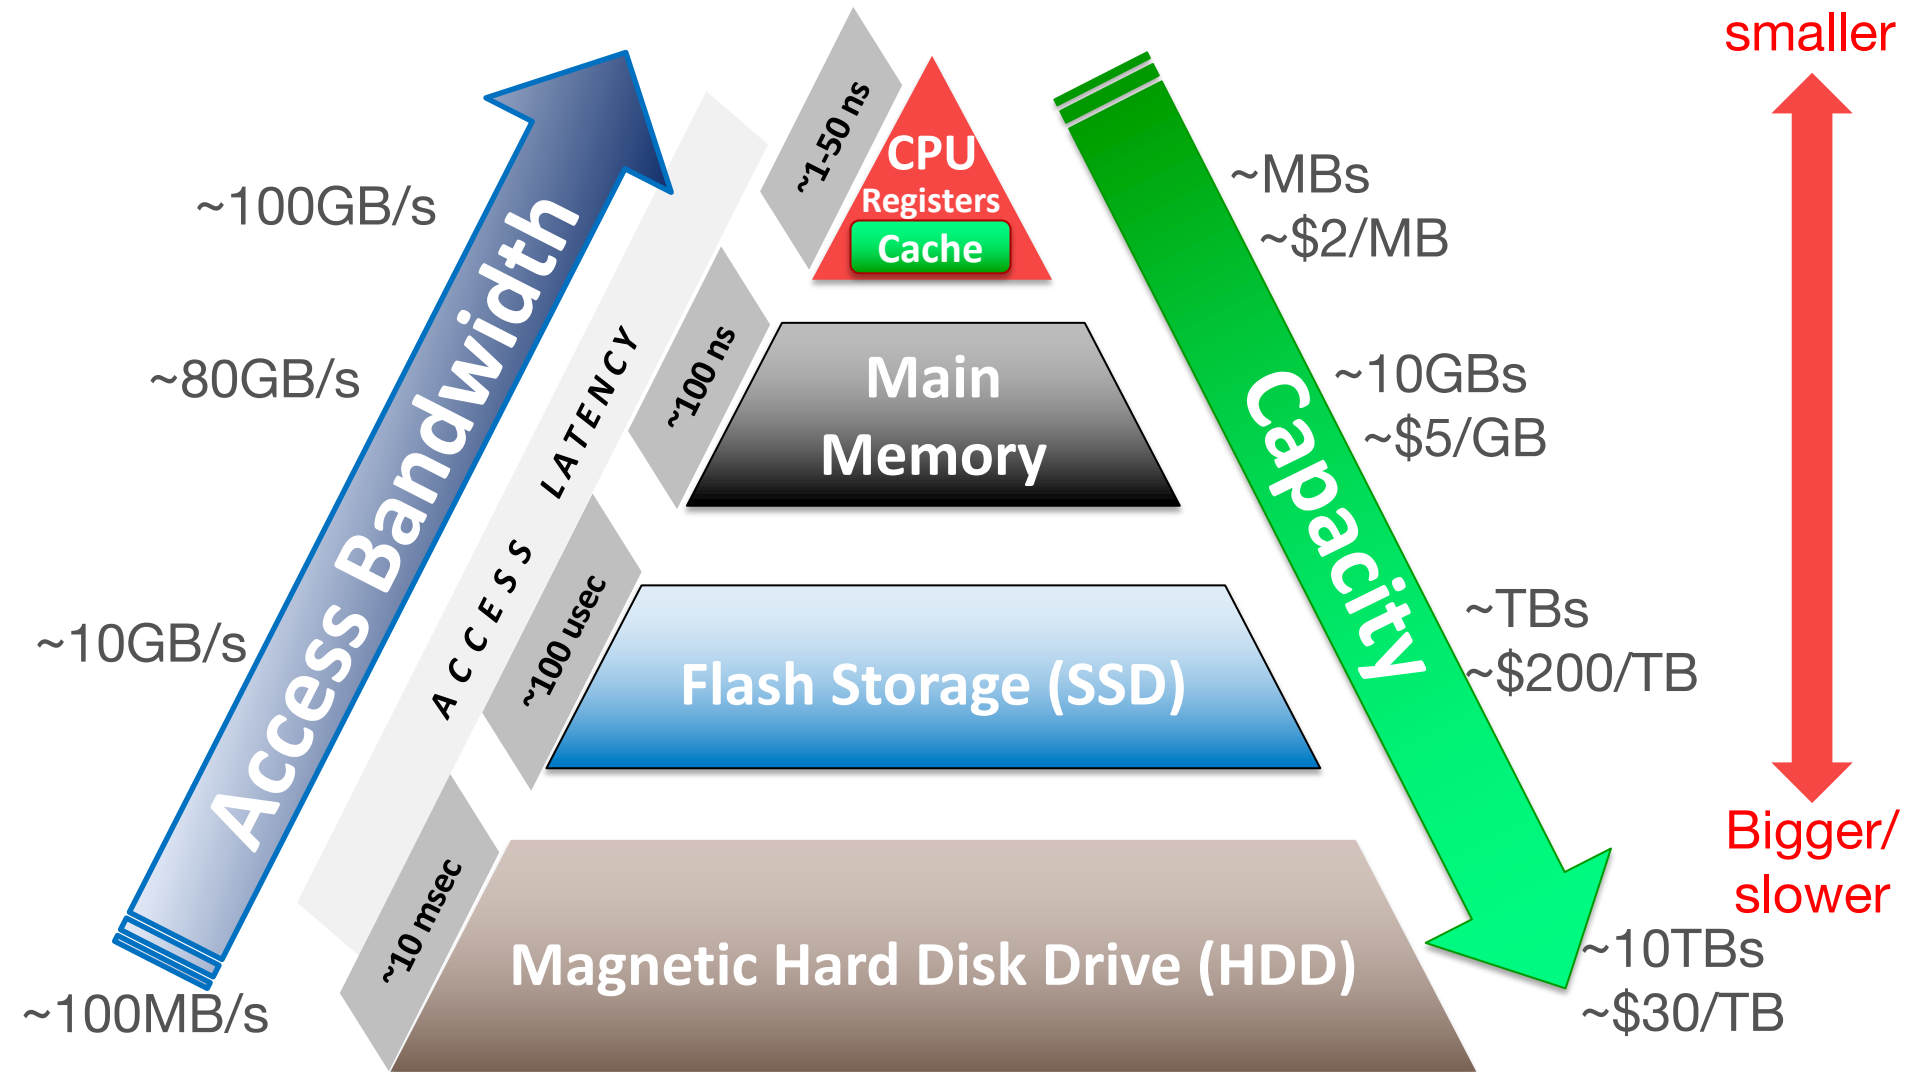
# Latency measurements

- Latency metrics
  - Average latency
  - Median latency
  - “Tail” latency (99<sup>th</sup> percentile, 99.9<sup>th</sup> percentile, etc.)

looking at the extreme occurrences of latency

- Which metrics do we expect **caching** to help with the most? average and sometimes median! Caching is a hit or miss so it can optimize the average latency (or sometimes median)


# Cache hierarchy



\*UCSD DSC 102: Systems for scalable analysis. Arun Kumar

# Resource tradeoffs

- File system caches file data in RAM
  - Uses **memory**
  - Avoids **storage** reads
- Browser caches recently visited pages as disk files
  - Uses **local storage** space
  - Avoids **network** transfers
- Python dictionary caches return values in a **dict** (key=args, val=return)
  - Uses **memory** space
  - Avoids **repeated compute**



```
cache = {}  
def f(x):  
    if not x in cache:  
        cache[x] = g(x)  
    return cache[x]
```

# Workload characteristics

## Application A

```
sum = 0
for i in range(0,1024):
    sum += a[i]
```

# Workload characteristics

## Application A

```
sum = 0
for i in range(0,1024):
    sum += a[i]
```

## Application B


```
import random

sum = 0
random.seed(1234);
for i in range(0,512):
    sum += a[random.randint(0,1023)]


random.seed(1234) # same seed
for i in range(0,512):
    sum += a[random.randint(0,1023)]
```

# Access patterns

Application A




Application B




# Access patterns

Application A



**Spatial Locality**

Application B



**Temporal Locality**



# Locality of data accesses

- **Spatial locality:**
  - Future access will be to nearby addresses
- **Temporal locality:**
  - Future access will be repeated to the same data

# Locality of data accesses

- **Spatial locality:**
  - Future access will be to nearby addresses
- **Temporal locality:**
  - Future access will be repeated to the same data
- Q: What is the implication of data locality to data systems applications?

# Locality optimization in Data Science

- Consider a matrix named `data` with  $16 \times 16$  elements
- Each row is of size 16 floats and **prefetching+caching** means 1/2 row of accessed data item is brought to CPU cache at a time

# Locality optimization in Data Science

- Consider a matrix named `data` with  $16 \times 16$  elements
- Each row is of size 16 floats and **prefetching+caching** means 1/2 row of accessed data item is brought to CPU cache at a time

- **Program 1**

```
for i in range(len(data[0])):  
    for row in data:  
        sum += row[i]
```

$16 \times 16 = 256$  CPU cache misses

Not too hardware-efficient (not able to exploit prefetching+caching)

# Locality optimization in Data Science

- Consider a matrix named `data` with  $16 \times 16$  elements
- Each row is of size 16 floats and **prefetching+caching** means 1/2 row of accessed data item is brought to CPU cache at a time

- **Program 1**

```
for i in range(len(data[0])):
    for row in data:
        sum += row[i]
```

$16 \times 16 = 256$  CPU cache misses

Not too hardware-efficient (not able to exploit prefetching+caching)

---

- **Program 2**

```
for row in data:
    for element in row:
        sum += element
```

Only  $16 \times 2$  CPU cache misses

- Each time  $\frac{1}{2}$  row of `data[i]` is prefetched to cache so subsequent accesses are hits!

# Peeking behind the scene...

- Data access granularity
  - If a process reads one byte and misses, **how much data should the CPU bring into the CPU cache?**
  - Tradeoff:
    - **Too little?** Will have many more misses if we read nearby bytes soon (recall spatial locality)
    - **Too much?** Wasteful to load data to cache that might never be accessed
- CPU caches data in units called **cache lines**
  - Typically, 64 bytes for modern CPUs (8 float64 numbers)

# Cache lines and misses

cache line

fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64

cache line

fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64

cache line

fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64

cache line

fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64

cache line

fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64

cache line

fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64  
fp64

How many misses?

How many misses?


How many misses?

# Memory layout of a matrix

Matrix of numbers  
**Logically**, 2-dimensional

Row
Row
Row
Row

**Physically**, those rows are arranged along **1-dimension** in the virtual address space



Virtual address space




# Memory layout of a matrix

Matrix of numbers  
**Logically**, 2-dimensional


Row
Row
Row
Row

Summing over row:  
data consolidated into a few cache lines (CPU cache friendly)

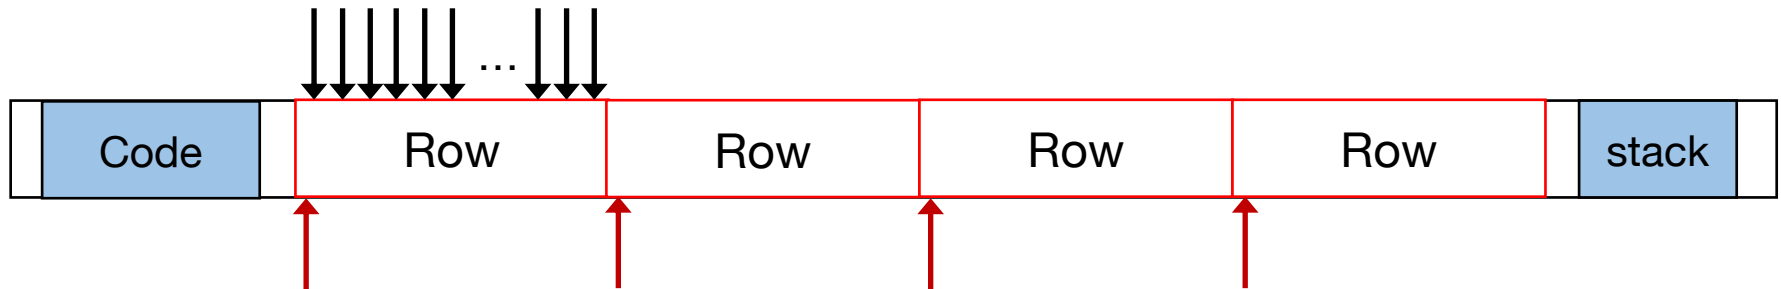


# Memory layout of a matrix

Matrix of numbers  
**Logically**, 2-dimensional



Summing over row:  
data consolidated into a few cache lines (CPU cache friendly)



Summing over column: each number is in its own cache line and triggers a **cache miss**

# Demo ...

# Caching policies

- When to **load** data to a cache?
  - Whenever the program reads something, add it to cache
- When to **evict** data from a cache (eviction policy)? Several policies:
  - Random: select any data at random for eviction
  - **FIFO** (first-in, first-out): evict whichever data that has been in the cache the longest
  - **LRU** (least recently used): evict which data that has been used the least recently

# Worksheet ...