

## Info

**Name:** Jacqui Unciano **Date:** 1 Feb. 2024 **Assignment:** Homework 3

### Follow this pattern:

Create a new notebook for your work. Parse the Frankenstein text to generate TOKENS and VOCAB tables.

```
In [1]: import pandas as pd
import numpy as np
import configparser
```

```
config = configparser.ConfigParser()
config.read('../../env.ini')
data_home = config['DEFAULT']['data_home']
output_dir = config['DEFAULT']['output_dir']
```

```
In [2]: text_file = f'{data_home}/gutenberg/pg42324.txt'
```

```
In [3]: clip_pats = [
    r"\*\*\s*START OF (? :THE|THIS) PROJECT",
    r"\*\*\s*END OF (? :THE|THIS) PROJECT"
]
chap_pat = r"(\s*(PREFACE\.)\s+)|(\s*(CHAPTER|LETTER)\s+X{0,3}(IX|IV|V?I{0,3})\s+)
```

```
In [4]: import sys
local_lib = config['DEFAULT']['local_lib']
sys.path.append(local_lib)
from textimporter import TextImporter
```

```
In [5]: my_text = TextImporter(src_file=text_file, ohco_pats=[('chap', chap_pat, 'm')], cli
my_text.import_source()
my_text.parse_tokens()
my_text.extract_vocab()
```

Importing /Users/jacqu/OneDrive/Documents/MSDS-at-UVA-2023/DS5001/data/gutenberg/pg42324.txt

Clipping text

Parsing OHCO level 0 chap\_id by milestone (\s\*(PREFACE\.)\s+)|(\s\*(CHAPTER|LETTER)\s+X{0,3}(IX|IV|V?I{0,3})\s+)

Parsing OHCO level 1 para\_num by delimiter \n\n

Parsing OHCO level 2 sent\_num by delimiter [.?!;:]+

Parsing OHCO level 3 token\_num by delimiter [\s',-]+

C:\Users\jacqu\OneDrive\Documents\MSDS-at-UVA-2023\DS5001\repo\lessons\lib\textimporter.py:117: UserWarning: This pattern is interpreted as a regular expression, and has match groups. To actually get the groups, use str.extract.

div\_lines = df[src\_col].str.contains(div\_pat, regex=True, case=False) # May want to parametrize case

```
Out[5]: <textimporter.TextImporter at 0x1ab5729aa10>
```

```
In [6]: TOKEN = my_text.TOKENS
        TOKEN.head()
```

```
Out[6]:
```

				token_str	term_str
chap_id	para_num	sent_num	token_num		
1	0	0	0	_To	to
			1	Mrs	mrs
		1	1	Saville	saville
			2	England	england
		2	0	_	

```
In [7]: my_text.VOCAB.head()
```

```
Out[7]:
```

	n	n_chars	p	s	i	h
term_str						
the	4200	3	0.055424	18.042857	4.173356	0.231302
and	2976	3	0.039272	25.463710	4.670371	0.183413
i	2854	1	0.037662	26.552207	4.730760	0.178168
of	2650	2	0.034970	28.596226	4.837753	0.169175
to	2105	2	0.027778	36.000000	5.169925	0.143609

Create a list of sentences from the TOKENS table and a list of terms from the VOCAB table.

```
In [8]: SENTS = my_text.gather_tokens(level=2)
        sents_list = SENTS.sent_num_str.tolist()
        sents_list[:5]
```

```
Out[8]: ['to mrs', 'saville england', '', 'st', 'petersburgh dec']
```

```
In [9]: term_list = my_text.VOCAB.index.tolist()
        term_list[:5]
```

```
Out[9]: ['the', 'and', 'i', 'of', 'to']
```

Generate ngram type tables and models, going up to the trigram level.

```
In [10]: def get_ngrams(TOKEN, n=2, sent_key='sent_num'):

        OHCO = TOKEN.index.names
        grouper = list(OHCO)[:OHCO.index(sent_key)+1]

        PADDED = TOKEN.groupby(grouper)\
            .apply(lambda x: '<s>' + ' '.join(x.term_str) + '</s>')\

```

```

        .apply(lambda x: pd.Series(x.split()))\
        .stack().to_frame('term_str')
    PADDED.index.names = grouper + ['token_num']

    NGRAMS = PADDED.groupby(grouper)\
        .apply(lambda x: pd.concat([x.shift(0-i) for i in range(n)], axis=1)).reset
    NGRAMS.index = PADDED.index
    NGRAMS.columns = [f'w{j}' for j in range(n)]

    return NGRAMS

```

```

In [11]: ngrams = 3
        widx = [f"w{i}" for i in range(ngrams)]

```

```

In [12]: def ngrams_to_models(ngrams):
        global widx
        n = len(ngrams.columns)
        model = [None for i in range(n)]
        for i in range(n):
            if i == 0:
                model[i] = ngrams.value_counts('w0').to_frame('n')
                model[i]['p'] = model[i].n / model[i].n.sum()
                model[i]['i'] = np.log2(1/model[i].p)
            else:
                model[i] = ngrams.value_counts(widx[:i+1]).to_frame('n')
                model[i]['cp'] = model[i].n / model[i-1].n
                model[i]['i'] = np.log2(1/model[i].cp)
        return model

```

```

In [13]: NG3 = get_ngrams(TOKEN, n=3)
        NG3.loc[(1,0,0)]

```

```

Out[13]:

```

	w0	w1	w2
token_num			
0	<s>	to	mrs
1	to	mrs	</s>
2	mrs	</s>	None
3	</s>	None	None

```

In [14]: M3 = ngrams_to_models(NG3)
        uni = M3[0].sort_values('n')
        bi = M3[1].sort_values('n')
        tri = M3[2].sort_values('n')

```

```

In [15]: uni

```

Out[15]:

	n	p	i
w0			
<b>irresolution</b>	1	0.000012	16.387580
<b>termination</b>	1	0.000012	16.387580
<b>brute</b>	1	0.000012	16.387580
<b>thonon</b>	1	0.000012	16.387580
<b>bruised</b>	1	0.000012	16.387580
...	...	...	...
<b>i</b>	2854	0.033289	4.908810
<b>and</b>	2976	0.034712	4.848421
<b>the</b>	4200	0.048989	4.351406
<b>&lt;/s&gt;</b>	5153	0.060105	4.056383
<b>&lt;s&gt;</b>	5153	0.060105	4.056383

6979 rows × 3 columns

In [16]: **bi**

Out[16]:

		n	cp	i
w0	w1			
<b>which</b>	<b>shot</b>	1	0.001792	9.124121
<b>hideous</b>	<b>that</b>	1	0.090909	3.459432
	<b>than</b>	1	0.090909	3.459432
	<b>narration</b>	1	0.090909	3.459432
	<b>monster</b>	1	0.090909	3.459432
...	...	...	...	...
<b>&lt;s&gt;</b>	<b>the</b>	366	0.071027	3.815497
	<b>and</b>	418	0.081118	3.623838
	<b>but</b>	457	0.088686	3.495147
<b>of</b>	<b>the</b>	530	0.200000	2.321928
<b>&lt;s&gt;</b>	<b>i</b>	820	0.159131	2.651717

40841 rows × 3 columns

In [17]: tri

Out[17]:

	w0	w1	w2	n	cp	i
	the	most	learned	1	0.017857	5.807355
	a	vast	and	1	0.333333	1.584963
			portion	1	0.333333	1.584963
			sheet	1	0.333333	1.584963
		vehicle	</s>	1	1.000000	0.000000
	...	...	...	...	...	...
	<s>	i	had	50	0.060976	4.035624
		and	i	62	0.148325	2.753163
		i	was	63	0.076829	3.702200
		it	was	67	0.515385	0.956279
		but	i	101	0.221007	2.177839

64838 rows × 3 columns

Write the code to answer the following questions:

## Question 1

List six words that precede the word "monster," excluding stop words (and sentence boundary markers). Stop words include 'a', 'an', 'the', 'this', 'that', etc. Hint: use the `df.query()` method.

In [18]: `bi.query('w1 == "monster").sort_values('n', ascending=False)`

Out[18]:

		n	cp	i
w0	w1			
the	monster	20	0.004762	7.714246
a	monster	3	0.002160	8.854868
hideous	monster	1	0.090909	3.459432
hellish	monster	1	0.142857	2.807355
detestable	monster	1	0.500000	1.000000
gigantic	monster	1	0.166667	2.584963
this	monster	1	0.002488	8.651052
miserable	monster	1	0.015385	6.022368
abhorred	monster	1	0.083333	3.584963
<s>	monster	1	0.000194	12.331197

Six words that precede the word, 'monster', are:

1. miserable
2. abhorred
3. gigantic
4. hideous
5. hellish
6. detestable

## Question 2

List the following sentences in ascending order of bigram perplexity according to the language model generated from the text:

The monster is on the ice.

Flowers are happy things.

I have never seen the aurora borealis.

He never knew the love of a family.

```
In [19]: ngrams = 2
         widx = [f"w{i}" for i in range(ngrams)]
```

```
In [20]: def sentence_to_token(sent_list, file=True):

         # Convert list of sentences to dataframe
         if file:
```

```

    S = pd.read_csv("test_sentences.txt", header=None, names=['sent_str'])
else:
    S = pd.DataFrame(sent_list, columns=['sent_str'])
S.index.name = 'sent_num'

# Convert dataframe of sentences to TOKEN with normalized terms
K = S.sent_str.apply(lambda x: pd.Series(x.split()).stack().to_frame('token_st
K['term_str'] = K.token_str.str.replace(r"[\W_]+", "", regex=True).str.lower()
K.index.names = ['sent_num', 'token_num']

return S, K

```

```

In [21]: test_sentences = """
The monster is on the ice
Flowers are happy things
I have never seen the aurora borealis
He never knew the love of a family
""".split("\n")[1:-1]
test_sentences

```

```

Out[21]: ['The monster is on the ice',
'Flowers are happy things',
'I have never seen the aurora borealis',
'He never knew the love of a family']

```

```

In [22]: TEST_SENTS, TEST_TOKENS = sentence_to_token(test_sentences, file=False)

```

```

In [23]: TEST_NGRAMS = get_ngrams(TEST_TOKENS)
TEST_NGRAMS.loc[0]

```

```

Out[23]:

```

	w0	w1
<b>token_num</b>		
0	<s>	the
1	the	monster
2	monster	is
3	is	on
4	on	the
5	the	ice
6	ice	</s>
7	</s>	None

```

In [24]: def test_model(model, ngrams, sents):

    global widx

    assert len(model) == len(ngrams.columns)

```

```

n = len(model)
ohco = ngrams.index.names

R = []
for i in range(n):
    T = ngrams.merge(model[i], on=widx[:i+1], how='left')
    T.index = ngrams.index
    T = T.reset_index().set_index(ohco + widx).i #.to_frame(f"i{i}")

    # This how we handle unseen combos
    T[T.isna()] = T.max()
    R.append(T.to_frame(f"i{i}"))

return pd.concat(R, axis=1)

```

```

In [25]: NG2 = get_ngrams(TOKEN)
        M = ngrams_to_models(NG2)

```

```

In [26]: R = test_model(M, TEST_NGRAMS, TEST_SENTS)

```

```

In [27]: def compute_perplexity(results, test_sents, n=2):
        for i in range(n):
            test_sents[f"pp{i}"] = np.exp2(results.groupby('sent_num')[f"i{i}"].mean())
        return test_sents

```

```

In [28]: PP = compute_perplexity(R, TEST_SENTS)
        PP.sort_values("pp1", ascending=True)

```

```

Out[28]:

```

	sent_str	pp0	pp1
sent_num			
0	The monster is on the ice	116.172431	80.655838
3	He never knew the love of a family	170.898523	136.954650
2	I have never seen the aurora borealis	341.056962	138.788045
1	Flowers are happy things	587.334984	534.302604

### Question 3

Using the bigram model represented as a matrix, explore the relationship between bigram pairs using the following lists. Hint: use the `.unstack()` method on the feature `n` and then use `.loc[]` to select the first list from the index, and the second list from the columns.

1. ['he','she'] to select the indices.
2. ['said','heard'] to select the columns.

```

In [46]: q3 = M[1].sort_values('n')
        q3 = q3.cp.unstack(fill_value=0)

```



```
In [47]: q3.loc([[ 'he', 'she'], [ 'said', 'heard']])
```

```
Out[47]:
```

	w1	said	heard
w0			
he	0.034483	0.008210	
she	0.011765	0.011765	

## Question 4

Generate 20 sentences using the .generate\_text() method from the langmod.NgramLanguageModel class.

```
In [72]: local_lib = config['DEFAULT']['local_lib']
sys.path.append(local_lib)
from langmod_class import NgramCounter
from langmod_class import NgramLanguageModel
```

```
In [73]: bigram = NgramCounter(sents=sents_list, vocab=term_list, n=2)
bigram.generate()
```

```
In [80]: bigram.S
```

```
Out[80]:
```

	sent_str	len
0	to mrs	4
1	saville england	4
2		2
3	st	3
4	petersburgh dec	4
...	...	...
5148	line 2863	4
5149	i do no not fear to die to i do now not fear t...	19
5150	fulfil the wishes of you parents to your parents	11
5151	end of the project gutenber ebook of frankens...	13
5152	shelley	3

5153 rows × 2 columns

```
In [83]: bimod = NgramLanguageModel(bigram)
```

```
In [84]: bimod.apply_smoothing()
```

```
In [86]: bimod.generate_text()
```

01. AS IF HER VOICE WOULD FOR WHEN I FEEL YOUR EYES AND HE SPOKE AND TRIUMPH AND SOMETIMES I HAD NO TRACE ITS CONCLUSION BY THE SUMMIT OF MY FAMILY IN STONE THAT OF MY WORK KNELT AT LENGTH HE THEN THOUGHT IT NECESSARY KNOWLEDGE OF MY POWER SHE COULD DRINK IN HORROR OF WHICH WAS MELANCHOLY SUBJECTS MY YOUTH BUT SAID THAT DAY MY DUTIES ON EXAMINING AND EXCITED BY EVERY THING WAS DECEIVED ALAS BUT I HAVE BEEN PENSIVE I AM CONTENT IF IMPATIENT TO UNDERSTAND THEIR LABOURS AND DELIGHT FRANKENSTEIN HAD CAST MY FRIENDS WITH ITS INTELLECTUAL EYE OF SWITZERLAND APPEARED IN BONDS OF RETURNING ALAS I HAD FIRST LITTLE FOOD AND TALKS AS IF I SHOULD BE THE ANCIENT STUDIES YOU WELL IF THEREFORE IN A RESISTLESS AND ALLOWED TO IMITATE HER PROMISED THAT THAT POSSESSED BY THE DREADFUL MEANS GAINED ADDITIONAL LOVE WILLIAM SAYING THIS SHORT AND IMMACULATE BEINGS WHO COMMITTED AT THAT WHAT I HOPED TO PROGNOSTICATE PEACE I WAS TO BESTOW ANIMATION WHEN I HAVE A LEAGUE IN LISTENING TO CLAIM THE BEAUTY THE HOVEL TO ME THEREFORE TO RECORD HAVE UNKNOWN HOW SHE DID NOT ESCAPE YET I CONFESS TO SECURE HIM AND PERISH IN PEACE.

## Question 5

Compute the redundancy  $R$  for each of the  $n$ -gram models using the MLE of the joint probability of each  $n$ gram type. In other words, for each model, just use the `.mle` feature as  $p$  in computing  $H = -\sum (p(\text{ng}) \log_2(1/p(\text{ng})))$ . Does  $R$  increase, decrease, or remain the same as the choice of  $n$ -gram increases in length? Hint: Remember that  $R = 1 - (H/H_{\text{max}})$ , where  $H$  is the actual entropy of the model and  $H_{\text{max}}$  is its maximum entropy.

```
In [119... NG3 = get_ngrams(TOKEN, n=3)
M3 = ngrams_to_models(NG3)
```

```
In [123... p = M3[0].p
Hmax = np.log2(len(M[0].index))
uniH = sum(p*np.log2(1/p))
1-(uniH/Hmax)
```

```
Out[123... 0.3085090204282228
```

```
In [124... p = M3[1].n / M3[1].n.sum()
Hmax = np.log2(len(M[0].index)**2)
uniH = sum(p*np.log2(1/p))
1-(uniH/Hmax)
```

```
Out[124... 0.4466133279213331
```

```
In [125... p = M3[2].n / M3[2].n.sum()
Hmax = np.log2(len(M[0].index)**3)
uniH = sum(p*np.log2(1/p))
1-(uniH/Hmax)
```

```
Out[125... 0.6310755519475555
```

The redundancy increases as the choice of n-gram increases in length.