

Info

Name: Jacqui Unciano **Date:** Feb 9, 2024 **Assignment:** HW04

Directions:

In this week's homework, you will use NLTK to help tokenize and annotate a small corpus of George Eliot's novels to create an F3 level digital analytical edition from them.

Using this week's Lab notebook M04_01_Pipeline.ipynb as a guide, import and combine the novels contained in the course repo directory /data/gutenberg/eliot-set.

You should produce the following related dataframes:

```
In [1]: import pandas as pd
import numpy as np
import configparser
from glob import glob
import re
import nltk

config = configparser.ConfigParser()
config.read('../../env.ini')
data_home = config['DEFAULT']['data_home']
output_dir = config['DEFAULT']['output_dir']
```

```
In [2]: source_files = f'{data_home}/gutenberg/eliot-set'
OHCO = ['book_id', 'chap_num', 'para_num', 'sent_num', 'token_num']
```

```
In [3]: import sys
local_lib = config['DEFAULT']['local_lib']
sys.path.append(local_lib)
from textimporter import TextImporter
from textparser import TextParser
```

```
src_file = f'{data_home}/gutenberg/eliot-set/ELIOT_GEORGE_THE_MILL_ON_THE_FLOSS-pg6688.txt'
LINES = pd.DataFrame(open(src_file, 'r', encoding='utf-8-sig').readlines(), columns=['line_str'])
LINES.index.name = 'line_num'
LINES.line_str = LINES.line_str.str.replace(r'\n+', ' ', regex=True).str.strip()
LINES.sample(5)
clip_pats = [r"[*]*\s*START OF (?:(THE|THIS) PROJECT", r"[*]*\s*END OF (?:(THE|THIS) PROJECT" ]
line_a = 0
line_b = len(LINES)
try: pat_a = LINES.line_str.str.match(clip_pats[0])
line_a = LINES.loc[pat_a].index[0] + 1
except: print("pat_a not found")
try: pat_b = LINES.line_str.str.match(clip_pats[1])
line_b = LINES.loc[pat_b].index[0] - 1
except: print("pat_b not found")
LINES = LINES.loc[line_a : line_b]
LINES.head()
LINES.tail()
roman = '[IVXLCM]+'
chap_pat = rf'^\s*Chapter\s+{roman}\. $"
chap_lines = LINES.line_str.str.match(chap_pat, case=False)
LINES.loc[chap_lines, 'chap_num'] = [i+1 for i in range(LINES.loc[chap_lines].shape[0])]
LINES.loc[chap_lines].head()
```

A library LIB with the following metadata (and data) about each book:

1. The book_id, matching the first level of the index in the CORPUS.
2. The raw book title will be sufficient, i.e. with title and author combined.
3. The path of the source file.
4. The regex used to parse chapter milestones.

5. The length of the book (number of tokens).
6. The number of chapters in the book.

```
In [4]: source_file_list = sorted(glob(f"{source_files}/*.txt"))
```

```
In [5]: book_data = []
for source_file_path in source_file_list:
    book_id = int(source_file_path.split("\\")[-1].split("-")[-1].split(".")[0].replace(
    book_title = source_file_path.split("\\")[-1].split("-")[0]
    book_data.append((book_id, source_file_path, book_title))
```

```
In [6]: LIB = pd.DataFrame(book_data, columns=['book_id', 'source_file_path', 'raw_title'])\
        .set_index('book_id').sort_index()
LIB.head()
```

```
Out[6]:
```

| | source_file_path | raw_title |
|---------|---|------------------------------------|
| book_id | | |
| 145 | /Users/jacqu/OneDrive/Documents/MSDS-at-UVA-20... | ELIOT_GEORGE_MIDDLEMARCH |
| 507 | /Users/jacqu/OneDrive/Documents/MSDS-at-UVA-20... | ELIOT_GEORGE_ADAM_BEDE |
| 6688 | /Users/jacqu/OneDrive/Documents/MSDS-at-UVA-20... | ELIOT_GEORGE_THE_MILL_ON_THE_FLOSS |

```
In [7]: roman = '[IVXLCM]+'
ohco_pat_list = [
    (145, rf"^\s*((PRELUDE\.)|(CHAPTER\s+{roman}\.))$"),
    (507, rf"^\s*(Chapter\s+{roman}|Epilogue)$"),
    (6688, rf"^\s*Chapter\s+{roman}\. $")
]
clip_pats = [
    r"\*\*\s*START OF (:THE|THIS) PROJECT",
    r"\*\*\s*END OF (:THE|THIS) PROJECT"
]
```

```
In [8]: LIB['chap_regex'] = LIB.index.map(pd.Series({x[0]:x[1] for x in ohco_pat_list}))
LIB.head()
```

Out[8]:

| | source_file_path | raw_title |
|---------|---|------------------------------------|
| book_id | | |
| 145 | /Users/jacqu/OneDrive/Documents/MSDS-at-UVA-20... | ELIOT_GEORGE_MIDDLEMARCH |
| 507 | /Users/jacqu/OneDrive/Documents/MSDS-at-UVA-20... | ELIOT_GEORGE_ADAM_BEDE [IV: |
| 6688 | /Users/jacqu/OneDrive/Documents/MSDS-at-UVA-20... | ELIOT_GEORGE_THE_MILL_ON_THE_FLOSS |

```
text = TextParser(src_file = LIB.iloc[0].source_file_path, ohco_pats=[('chap', LIB.iloc[0].chap_regex, 'm')], clip_pats = clip_pats,
use_nltk=True) text.import_source() text.parse_tokens() text.extract_vocab() TOKEN = text.TOKENS TOKEN.sample(10)
```

```
In [9]: num_tokens = []
num_chaps = []
blist = []

for i in range(len(LIB)):
    text = TextParser(src_file = LIB.iloc[i].source_file_path,
                      ohco_pats=[('chap', LIB.iloc[i].chap_regex, 'm')],
                      clip_pats = clip_pats,
                      use_nltk=True)
    text.import_source()
    text.parse_tokens()
    text.extract_vocab()
    TOKEN = text.TOKENS
    CHAPS = text.gather_tokens(level=0)
    num_tokens.append((LIB.index[i], len(TOKEN)))
    num_chaps.append((LIB.index[i], len(CHAPS)))
    blist.append((LIB.index[i], TOKEN))

LIB['num_tokens'] = LIB.index.map(pd.Series({x[0]:x[1] for x in num_tokens}))
LIB['num_chaps'] = LIB.index.map(pd.Series({x[0]:x[1] for x in num_chaps}))
```

```
C:\Users\jacqu\OneDrive\Documents\MSDS-at-UVA-2023\DS5001\repo\lessons\lib\textparse
r.py:132: UserWarning: This pattern is interpreted as a regular expression, and has
match groups. To actually get the groups, use str.extract.
```

```
div_lines = self.TOKENS[src_col].str.contains(div_pat, regex=True, case=True) # TO
DO: Parametrize case
```

```
line_str chap_str
Index(['chap_str'], dtype='object')
```

```
C:\Users\jacqu\OneDrive\Documents\MSDS-at-UVA-2023\DS5001\repo\lessons\lib\textparse
r.py:132: UserWarning: This pattern is interpreted as a regular expression, and has
match groups. To actually get the groups, use str.extract.
```

```
div_lines = self.TOKENS[src_col].str.contains(div_pat, regex=True, case=True) # TO
DO: Parametrize case
```

```
line_str chap_str
Index(['chap_str'], dtype='object')
line_str chap_str
Index(['chap_str'], dtype='object')
```

In [10]: `LIB.head()`

Out[10]:

| | source_file_path | raw_title |
|---------|---|------------------------------------|
| book_id | | |
| 145 | /Users/jacqu/OneDrive/Documents/MSDS-at-UVA-20... | ELIOT_GEORGE_MIDDLEMARCH |
| 507 | /Users/jacqu/OneDrive/Documents/MSDS-at-UVA-20... | ELIOT_GEORGE_ADAM_BEDE [IV: |
| 6688 | /Users/jacqu/OneDrive/Documents/MSDS-at-UVA-20... | ELIOT_GEORGE_THE_MILL_ON_THE_FLOSS |

A an aggregate of all the novels' tokens CORPUS with an appropriate OHCO index, with following features:

1. The token string.
2. The term string.
3. The part-of-speech tag inferred by NLTK.TK.

In [11]: `CORPUS = pd.concat([i[1] for i in blist], keys=[i[0] for i in blist])`
`CORPUS.index.names = OHCO[:]`

In [12]: `CORPUS.sample(10)`

Out[12]:

| | | | | | pos_tuple | pos | token_str | term_str |
|---------|----------|----------|----------|-----------|--------------|-----|-----------|----------|
| book_id | chap_num | para_num | sent_num | token_num | | | | |
| 6688 | 25 | 58 | 2 | 9 | (s, NN) | NN | s | s |
| 507 | 52 | 4 | 3 | 7 | (thee—I, NN) | NN | thee—I | theei |
| 145 | 69 | 18 | 3 | 16 | (., .) | . | . | NaN |
| 507 | 53 | 24 | 2 | 6 | (be, VB) | VB | be | be |
| | 18 | 11 | 6 | 7 | (to, TO) | TO | to | to |
| 145 | 55 | 38 | 0 | 33 | (and, CC) | CC | and | and |
| | 35 | 4 | 9 | 6 | (being, VBG) | VBG | being | being |
| 507 | 45 | 2 | 6 | 65 | (,, ,) | , | , | NaN |
| 6688 | 14 | 9 | 7 | 9 | (by, IN) | IN | by | by |
| | 58 | 54 | 3 | 19 | (,, ,) | , | , | NaN |

A vocabulary VOCAB of terms extracted from CORPUS, with the following annotation features derived from either NLTK or by using operations presented in the notebook:

1. Stopwords.
2. Porter stems.
3. Maximum POS; i.e. the most frequently associated POS tag for the term using .idxmax().
Note that ties are handled by the method.
4. POS ambiguity expressed a number of POS tags associated with a term's tokens.ens.

```
In [13]: VOCAB = CORPUS.groupby(level=0)['term_str'].value_counts().to_frame('n')
VOCAB.index.names = ['book_id', 'term_str']
# VOCAB.head()
```

```
In [14]: sw = pd.DataFrame(nltk.corpus.stopwords.words('english'), columns=['term_str'])
sw = sw.reset_index().set_index('term_str')
sw.columns = ['dummy']
sw.dummy = 1
```

```
sw.head()
```

```
In [15]: VOCAB['stop'] = VOCAB.index.get_level_values(1).map(sw.dummy)
VOCAB['stop'] = VOCAB['stop'].fillna(0).astype('int')
```

```
In [16]: VOCAB[VOCAB.stop == 1].sample(10)
```

Out[16]:

| | | n | stop |
|---------|----------|-----|------|
| book_id | term_str | | |
| 507 | they | 750 | 1 |
| | our | 255 | 1 |
| 6688 | only | 354 | 1 |
| | been | 732 | 1 |
| | needn | 11 | 1 |
| | after | 226 | 1 |
| 507 | ours | 5 | 1 |
| | through | 143 | 1 |
| 6688 | such | 214 | 1 |
| 507 | only | 315 | 1 |

```
In [17]: from nltk.stem.porter import PorterStemmer
stemmer = PorterStemmer()
VOCAB['stem_porter'] = VOCAB.index.get_level_values(1).map(lambda x: stemmer.stem(x))

from nltk.stem.snowball import SnowballStemmer
stemmer2 = SnowballStemmer("english")
VOCAB['stem_snowball'] = VOCAB.index.get_level_values(1).map(lambda x: stemmer2.stem(x))

from nltk.stem.lancaster import LancasterStemmer
stemmer3 = LancasterStemmer()
VOCAB['stem_lancaster'] = VOCAB.index.get_level_values(1).map(lambda x: stemmer3.stem(x))

VOCAB.sample(10)
```

Out[17]:

| | | n | stop | stem_porter | stem_snowball | stem_lancaster |
|---------|------------|----|------|-------------|---------------|----------------|
| book_id | term_str | | | | | |
| 6688 | agreeable | 25 | 0 | agreeabl | agreeabl | agr |
| | dictated | 2 | 0 | dictat | dictat | dict |
| 507 | sky | 19 | 0 | sky | sky | sky |
| 145 | slips | 2 | 0 | slip | slip | slip |
| 6688 | unmapped | 1 | 0 | unmap | unmap | unmap |
| 507 | convulsion | 1 | 0 | convuls | convuls | convuls |
| 145 | casting | 3 | 0 | cast | cast | cast |
| | bedrooms | 1 | 0 | bedroom | bedroom | bedroom |
| | fling | 4 | 0 | fling | fling | fling |
| | tickled | 1 | 0 | tickl | tickl | tickl |

```
In [18]: VOCAB['max_pos'] = CORPUS.groupby(level=0)[['term_str', 'pos']].value_counts().unsta
```

```
In [19]: VOCAB['n_pos'] = CORPUS.groupby(level=0)[['term_str', 'pos']].value_counts().unstack
VOCAB['cat_pos'] = CORPUS.groupby(level=0)[['term_str', 'pos']].value_counts().to_fr
        .groupby(['book_id', 'term_str']).pos.apply(lambda x: set(x))
```

```
In [20]: VOCAB.sample(10)
```

Out[20]:

| | | n | stop | stem_porter | stem_snowball | stem_lancaster | max_pos |
|---------|---------------|---|------|---------------|---------------|----------------|---------|
| book_id | term_str | | | | | | |
| 507 | footsteps | 4 | 0 | footstep | footstep | footstep | NNS |
| | unlocked | 2 | 0 | unlock | unlock | unlock | VBD |
| 6688 | corpses | 2 | 0 | corps | corps | corps | NNS |
| 507 | insist | 2 | 0 | insist | insist | insist | NN |
| 6688 | godmother | 3 | 0 | godmoth | godmoth | godmoth | NN |
| 145 | mothers | 7 | 0 | mother | mother | moth | NNS |
| 507 | goodfornought | 1 | 0 | goodfornought | goodfornought | goodfornought | JJ |
| | comforted | 9 | 0 | comfort | comfort | comfort | VBN |
| 6688 | wrapped | 4 | 0 | wrap | wrap | wrap | VBD |
| 145 | oblivion | 1 | 0 | oblivion | oblivion | obl | NN |

Once you have these, use the dataframes to answer these questions:

Question 1: What regular expression did you use to chunk *Middlemarch* into chapters?

```
In [21]: LIB.chap_regex.loc[145]
```

```
Out[21]: '^\\s*((PRELUDE\\.\\.)|(CHAPTER\\s+[IVXLCM]+\\.\\.))$'
```

Question 2: What is the title of the book that has the most tokens?

```
In [22]: LIB.loc[LIB['num_tokens'].idxmax()].raw_title
```

```
Out[22]: 'ELIOT_GEORGE_MIDDLEMARCH'
```

Question 3: How many chapter level chunks are there in this novel?

```
In [23]: LIB.loc[LIB['num_tokens'].idxmax()].num_chaps
```

```
Out[23]: 87
```

Question 4: Among the three stemming algorithms -- Porter, Lancaster, and Snowball -- which is the most aggressive, in terms of the number of words associated with each stem?

```
In [24]: p = VOCAB.stem_porter.value_counts().to_frame()
p.sort_values(by='count', ascending=False).head()
```


Out[24]:

| | count |
|-------------|-------|
| stem_porter | |
| admir | 26 |
| observ | 26 |
| respect | 24 |
| impress | 24 |
| continu | 23 |

```
In [25]: s = VOCAB.stem_snowball.value_counts().to_frame()
s.sort_values(by='count', ascending=False).head()
```

Out[25]:

| | count |
|---------------|-------|
| stem_snowball | |
| admir | 29 |
| observ | 26 |
| wonder | 25 |
| impress | 24 |
| respect | 24 |

```
In [26]: l = VOCAB.stem_lancaster.value_counts().to_frame()
l.sort_values(by='count', ascending=False).head()
```

Out[26]:

| | count |
|----------------|-------|
| stem_lancaster | |
| cont | 74 |
| man | 47 |
| adv | 44 |
| pass | 43 |
| not | 39 |

Answer 4: I would say that Lancaster was the most aggressive since there are a lot of terms that have been chopped into their supposed stem. It really goes at it with stemming the words.

Question 5: Using the most aggressive stemmer from the previous question, what is the stem with the most associated terms?

```
In [27]: l.sort_values(by='count', ascending=False).head(1)
```

```
Out[27]:
```

| | count |
|----------------|-------|
| stem_lancaster | |
| cont | 74 |

Answer 5: Using the Lancaster stemmer, the stem with the most associated terms is 'cont'.